CNN Baseline Report

A Convolutional Neural Network is designed to process the four digits and predict their sum. The architecture consists of:

- **Convolutional Layers**: Two layers with ReLU activation layers and max pooling to extract features from the input
- **Fully Connected Layers**: A flattening layer followed by two dense layers
- **Output Layer**: A single neuron outputting the predicted sum

Training Details:

- **Loss Function**: MSE loss is used to minimize the difference between predicted and actual sums
- **Learning Rate:** 0.001

The model is trained for 10 epochs. Preliminary results indicate that the CNN model cannot learn the sum without adequate thresholding/clustering of the 4 digits.